

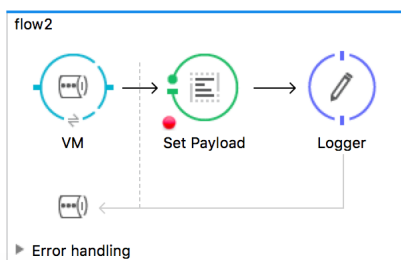
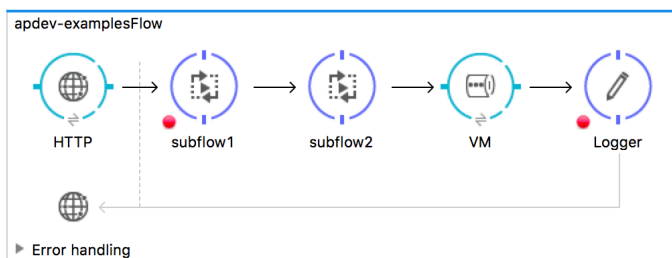
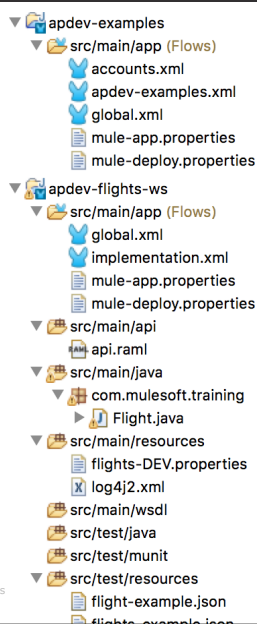




Module 6: Structuring Mule Applications

Goal



Objectives



- Create and reference flows and subflows
- Pass messages between flows using the Java Virtual Machine (VM) transport
- Investigate variable persistence through subflows and flows and across transport barriers
- Encapsulate global elements in a separate configuration file
- Explore the files and folder structure of a Mule project

All contents © MuleSoft Inc.

3

Encapsulating processors into separate flows and subflows



Break up flows into separate flows and subflows




- Makes the graphical view more intuitive
 - You don't want long flows that go off the screen
- Makes XML code easier to read
- Enables code reuse
- Provides separation between an interface and implementation
 - We already saw this
- Makes them easier to test

Subflows



- Subflows are executed exactly as if the processors were still in the calling flow
 - Always run synchronously in the same thread
 - Inherit the processing and exception strategies of the flow that triggered its execution



Flows



- Flows, on the other hand, have much more flexibility in how they are used
 - Can have their own processing and exception strategies
 - Can be synchronous or asynchronous
- Flows without message sources are sometimes called private flows

What is a flow processing strategy?




- A flow **processing strategy** determines how Mule implements message processing for a given flow
 - Should the message be processed synchronously (on the same thread) or asynchronously (on a different thread)?
 - If asynchronously, what are the properties of the pool of threads used to process the messages?
 - If asynchronously, how will messages wait for their turn to be processed in the second thread?

Flows contain 3 thread pools



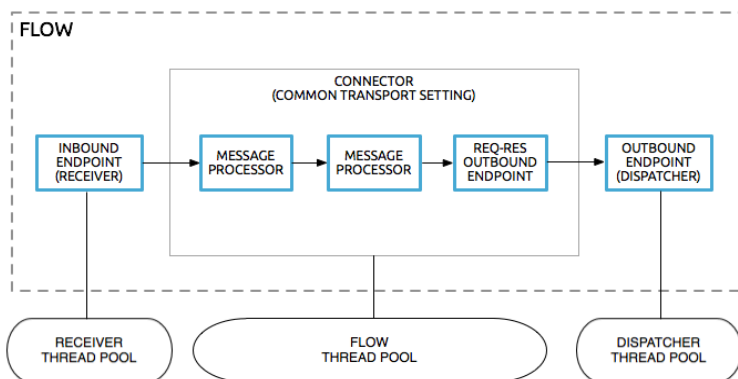
- Receiving
 - Message source's threads
- Flow processing
 - Message processor's threads
- Dispatching
 - Outbound endpoint's threads



All contents © MuleSoft Inc.

19

Usage of pools depends on a flow's behavior



Synchronous Flow

Asynchronous Flow (SEDA queue)

All contents © MuleSoft Inc.

20

Staged Event Driven Architecture (SEDA)



- The architecture upon which Mule was built
- Decouples receiving, processing, and dispatching phases
- Supports higher levels of parallelism in specific stages of processing
- Allows for more-specific tuning of areas within a flow's architecture

What determines a flow's processing strategy?




- Mule automatically sets a flow to be
 - Synchronous or queued-asynchronous
- A flow is set to synchronous if
 - The message source is request-response
 - The sender of the message is expecting a response
 - The flow partakes in a transaction
- Otherwise a flow is set to queued-asynchronous
 - The message source is not expecting a response

Synchronous flows




- When a flow receives a message, all processing, including the processing of the response, is done in the same thread
- Uses only the message source's thread pool
- The flow's thread pool is elastic and will have one idle thread that is never used
- Tuning for higher-throughput happens on the connector receiver's level



All contents © MuleSoft Inc.

25

Queued-asynchronous flows



- Decouples and uses all 3 thread pools
- Uses queues, whose threads drop messages off for the subsequent pool's thread to pickup
- Pools, queues, and behaviors of this strategy are configurable
- By default, the flow thread pool has 16 threads

All contents © MuleSoft Inc.

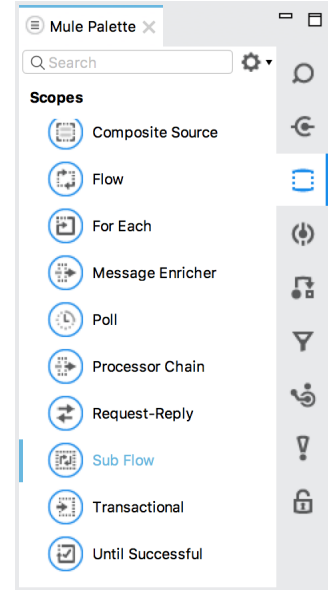
26

Creating flows and subflows



- Use Flow scope to create a new flow
 - Or drag any message processor to the canvas
- Use Sub flow scope to create subflows
- Use Flow Reference component to pass messages to other flows or subflows
- Flow variables persist through all flows unless the message crosses a transport boundary
 - We will explore this shortly


All contents © MuleSoft Inc.



Setting the flow processing strategy



- The flow processing strategy is automatically set
- It can be changed in the flow's properties view
- This is usually done when a custom queued-asynchronous profile has been created for tuning application performance
 - This is covered in the *Anypoint Platform Development: Advanced* course



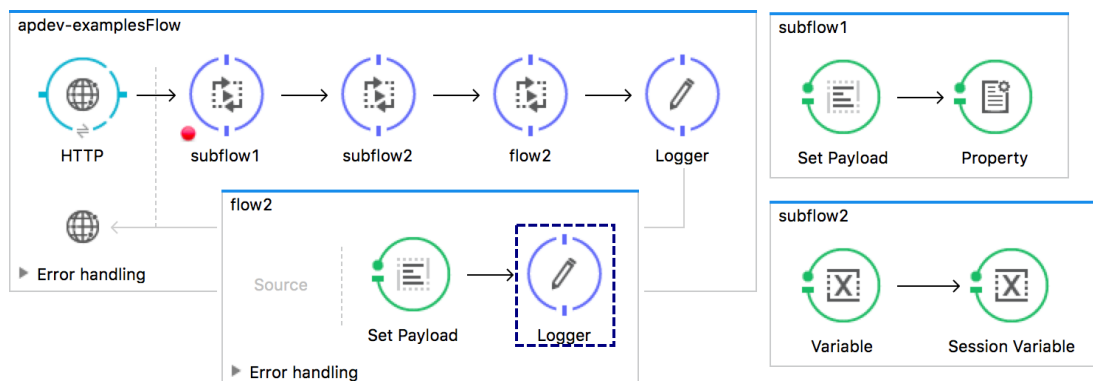
All contents © MuleSoft Inc.

27

Walkthrough 6-1: Create and reference flows and subflows




- Extract processors into separate subflows and flows
- Use the Flow Reference component to reference other flows
- Explore variable persistence through flows and subflows



28

Using VM endpoints in flows

The Java Virtual Machine (VM) transport



- The Java Virtual Machine (VM) transport can be used for intra-JVM communication between Mule flows
- Each app in a Mule instance has its own, unique set of VM endpoints
- The VM transport can only handle communications within an app or between apps in the same domain
- This transport by default uses in-memory queues but can optionally be configured to use persistent queues
 - VM file persistency does not work on clusters

All contents © MuleSoft Inc.

30

Using the VM transport



- Before Mule 3, the VM transport was needed to pass a message from one flow to another
- In Mule 3, Flow Reference was added to let flows directly reference one another without a transport in the middle
 - This obviates the need for VM in many cases
- VM transport is now mostly used to
 - Achieve higher levels of parallelism in specific stages of processing
 - Allow for more-specific tuning of areas within a flow's architecture
 - Call flows in other applications that are in the same domain
- Using VMs to tune application performance is covered in the *Anypoint Platform Development: Advanced* course

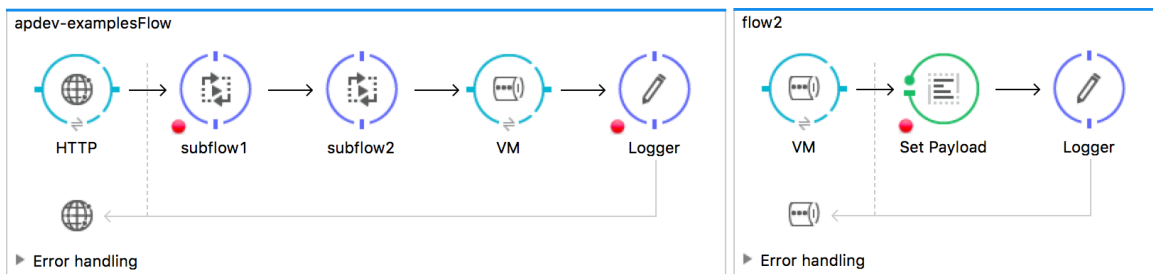
All contents © MuleSoft Inc.

31

Walkthrough 6-2: Pass messages between flows using the Java Virtual Machine (VM) transport



- Pass messages through an HTTP transport barrier
- Pass messages between flows using the VM transport
- Explore variable persistence across transport barriers



Organizing Mule application files

Separating apps into multiple configuration files



- Just as we separated flows into multiple flows, we also want to separate configuration files into multiple configuration files
- Monolithic files are difficult to read and maintain
- Separating an application into multiple configuration files makes code
 - Easier to read
 - Easier to work with
 - Easier to test
 - More maintainable

All contents © MuleSoft Inc.

25

Encapsulating global elements in a configuration file



- It can be confusing if you reference global elements in one file that are defined in various, unrelated files
- It also makes it hard to find them
- A good solution is to put most global elements in one config file
 - All the rest of the files reference them
 - If a global element is specific to and only used in one file, it makes sense to keep it in that file

All contents © MuleSoft Inc.

26

Creating multiple applications



- You are also not going to want all your flows in one application/project
- Separate functionality into multiple applications to
 - Allow managing and monitoring of them as separate entities
 - Use different, incompatible JAR files
- Run more than one application at a time in Anypoint Studio by creating a run configuration

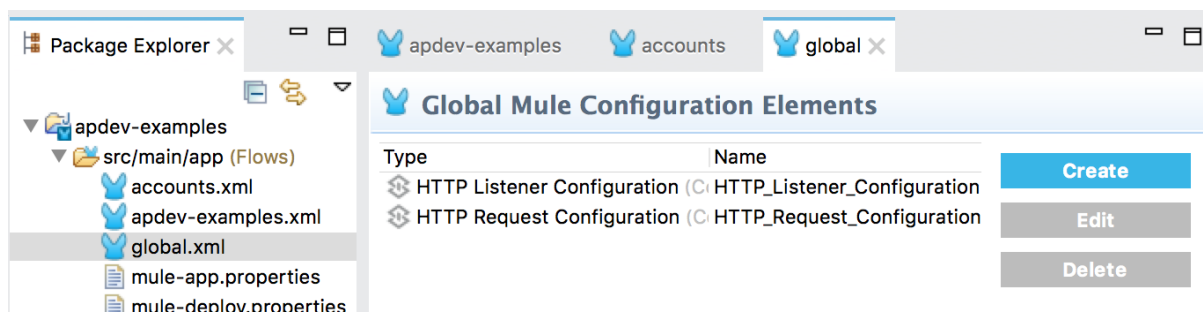
All contents © MuleSoft Inc.

27

Walkthrough 6-3: Encapsulate global elements in a separate configuration file



- Create a new configuration file with an endpoint that uses an existing global element
- Create a configuration file global.xml for just global elements
- Move the existing global elements to global.xml



All contents © MuleSoft Inc.

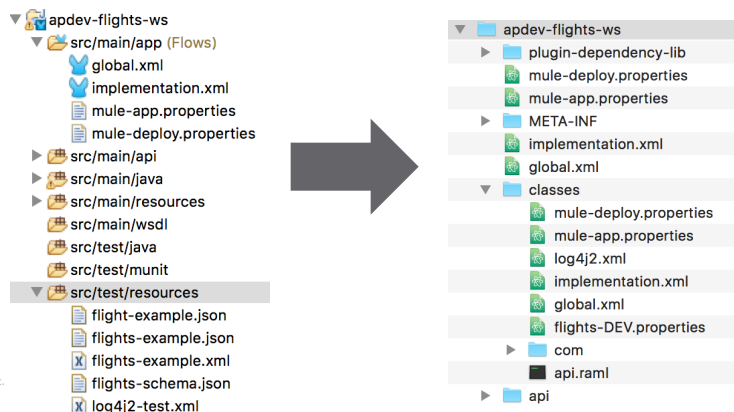
28

Organizing Mule project files

Examining the folder structure for a Mule project



- The names of folders indicate what they should contain
- src/test folders should contain files only needed at development time
 - Like schema and example files for metadata types, sample data for transformations
 - They are not included in the application ZIP when it is packaged

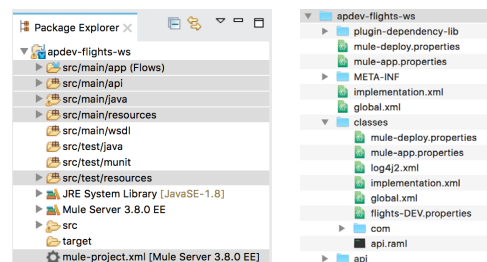


All contents © MuleSoft Inc.

Walkthrough 6-4: Create a well organized Mule project



- Create a project with a new RAML file that is added to Anypoint Platform
- Review the project's configuration and properties files
- Create an application properties file and a global configuration file
- Add Java files and test resource files to the project
- Create and examine the contents of a deployable archive for the project



All contents © MuleSoft Inc.

31

Sharing resources between applications



Sharing resources between applications



- **Shared domains** can be used to share some connectors and endpoints between applications
- Enables multiple development teams to work in parallel using the same set of reusable connectors
- Defining these connectors as shared resources at the domain level allows the team to
 - Expose multiple services within the domain on same port
 - Share the connection to persistent storage
 - Ensure consistency between applications upon any changes, as the configuration is only set in one place

All contents © MuleSoft Inc.

33

Using domains to share resources between apps



- Only available on-prem, not on CloudHub
- The general process is to
 - Create a Mule Domain Project and associate Mule applications with a domain
 - Define a set of resources (and the libraries required by those resources) for a domain to share between the apps
- Currently, only the following can be shared
 - HTTP/HTTPS (both endpoints and connectors), WMQ, Database, JMS, VM, JBoss and Bitronix Transaction Managers

All contents © MuleSoft Inc.

34

Sharing code between applications



- You may also want to share code between applications
 - Connector configurations
 - Endpoint definitions
 - Flows and subflows
 - Transformers
 - DataWeave DWL files
- To share code between applications, use **Maven and its dependency management mechanism**
- Maven is a software project management tool that helps
 - Control the build lifecycle of software projects
 - Manage project dependencies

All contents © MuleSoft Inc.

35

Using Maven to share code between applications



- Using Maven is covered in the *Anypoint Platform Development: Advanced* course
- The general process is to
 - Create a shared XML configuration file containing the code to share
 - Compile it into a JAR file
 - Use Maven's dependency mechanism to include it in your application
 - Import the configuration file for reference
 - <https://docs.mulesoft.com/mule-user-guide/v/3.8/sharing-custom-configuration-fragments>

All contents © MuleSoft Inc.

36

Summary



Summary



- Separate a flow into multiple flows for easier to read files and for reusability
- Extract duplicate code into reusable private flows or subflows
- Private flows are flows without message sources
- Subflows are executed exactly as if in the calling flow
 - Always run synchronously
 - Inherit processing and exception strategies of the calling flow
- A flow processing strategy determines how Mule implements message processing for a given flow
 - All Mule flows have an implicit processing strategy which Mule applies automatically
 - Endpoints with a request-response exchange pattern are set to synchronous
 - Endpoints with a one-way exchange pattern are set to queued-asynchronous

All contents © MuleSoft Inc.

38

Summary



- Use Flow Reference to let flows directly reference one another without a transport in the middle
- Use the Java Virtual Machine (VM) transport for intra-JVM communication between Mule flows
 - To achieve higher levels of parallelism in specific stages of processing
 - To allow for more-specific tuning of areas within a flow's architecture
 - To call flows in other applications that are in the same domain
- Flow variables persist through all flows unless the message crosses a transport boundary (like VM or HTTP)
- Session variables persist across some but not all transport barriers (VM yes, HTTP no)

All contents © MuleSoft Inc.

40

Summary



- Separate application functionality into multiple configuration files for easier development and maintenance
 - Encapsulate global elements that will be used in multiple configuration files into their own separate configuration file
- Put files only needed at development time in the src/test folders
 - They are not included in the application ZIP when it is packaged
- Separate functionality into multiple applications to allow managing and monitoring of them as separate entities
- Share resources between applications by creating a shared domain

All contents © MuleSoft Inc.

37